

# AutoEncoder

## Model

- We have an encoder network that compresses the input to a small vector
- We have a decoder network that decodes the vector back to an approximation of the input
- We minimise the loss between the output of the decoder network and the input
- Unfortunately the space of encoded vectors is a bit strange so it's hard to figure out what's going on
- You can add convolutional layers in – this is generally a pretty good move